UNIVERSITÄT OSNABRÜCK

Faculty of
Cognitive Science

Bachelor Thesis

Romeo: An Interactive 3D Visualization Library for Julia

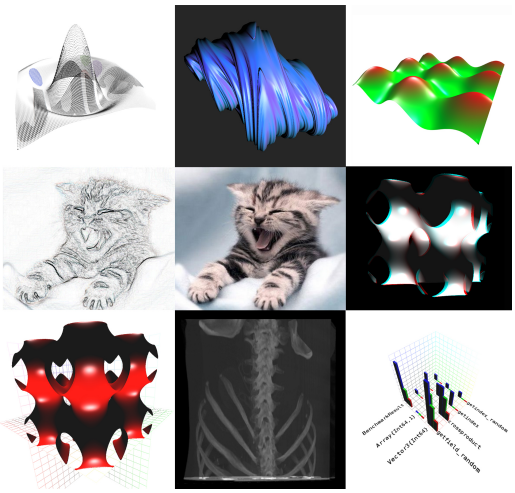

Author: Simon Danisch
sdanisch@email.de

Supervisor: Prof. Dr.-Ing. Elke Pulvermüller

Co-Reader: Apl. Prof. Dr. Kai-Christoph Hamborg

Filing Date: 01.02.2014

I Abstract

This bachelor thesis is about writing a simple scripting environment for scientific computing, with focus on visualizations and interaction. Focus on visualization means that every variable can be inspected and visualized at runtime, ranging from a textual representation to complex 3D scenes. Interaction is achieved by offering simple GUI elements for all parts of the program and the visualizations. All libraries are implemented in Julia and modern OpenGL, to offer high performance, opening the world to scientists who have to work with large datasets. Julia is a novel high-level programming language for scientific computing, promising to match C speed, making it the optimal match for this project. -This section needs more work, and should probably be written in the end

II Table of Contents

III List of Figures

IV List of Tables

V Listing-Verzeichnis

VI List of Abbreviations

GUI	Graphical User Interface
LLVM	Low Level Virtual Machine
IR	Intermediate Representation
gcc	GNU Compiler Collection
Matlab	Matrix Laboratory
REPL	Read Eval Print Loop
GPU	Graphics Processing Unit
GLSL	OpenGL Shading Language
OpenCL	Open Compute Language
OpenGL	Open Graphics Language
VTK	Visualization Toolkit
AST	Abstract Syntax Tree
CUDA	Compute Unified Device Architecture

1 Introduction

This Bachelor Thesis is about writing a fast and interactive 3D visualization environment for scientific computing. The name of the library is Romeo, but other libraries had to be developed in order to achieve the functionality. The focus is on usability, applied to all the different interfaces, ranging from abstract API interfaces to graphical user interfaces. The ultimate goal is to make scientific computing more accessible to the user. As Graphical User Interface (GUI) elements and editable text fields are supplied, one can also write and execute scripts. Using these widgets, all bound variables can be visualized and some of them can be edited interactively. This can be used as a basis for interactive programming or visual debugging, further helping the user to understand his algorithms.

The introduction is structured in the following way. First, an introduction to the general field of research and its challenges is given. From these challenges, the problems relevant to this thesis will be extracted. Finally this chapter will conclude with a solution to the problem, how to measure the success and give an outlook on the structure of the entire Bachelor Thesis.

1.1 Scientific Computing

Scientific computing is the area of computing that evolves around all kind of scientific research. It is a very broad field involving a lot of different challenges. In some areas like particle physics, the problems are computationally so demanding, that they can only be solved with the help of super computers. In other areas like robotics, it is important to be efficient, because the algorithms are running on embedded systems with limited resources. In a lot of other areas, speed does not need to be important, but it can be that the algorithm in itself is very difficult to comprehend. So the more comprehendable an algorithm can be written down in a programming language, the easier it will be to implement the algorithm without errors. Above all, programming itself is secondary to the research goal. This means that it can be expected that a researcher just has rudimentary programming skills. Even if he is a professional programmer, he wants to put as little time as possible into solving pure programming problems.

So things like manual memory management and difficult design patterns with a lot of boilerplates are to be avoided in scientific computing. This has lead to the rise of programming languages and tools specifically tailored to scientific computing. The most prominent examples include Mathematica, R, and Matlab. Python could be in this list as well, but the scientific computing part is only realized by third party libraries, while Python itself is a multi purpose language. The others all aim to provide simple syntax for linear algebra and statistical code, while taking away programmatically difficult tasks like

memory management. Also, they come with a rich standard library, which means most research can be done without loading any additional module, which makes them great tools for rapid prototyping. At the current state, the speed of these languages suffer from the high level of abstraction. In order to cater to the fields of scientific computing which is in need of highly performant code, a lot of the core is written in another language like C/C++ and Fortran. This poses a problem in itself. As soon as a researcher needs to do something out of the ordinary in a performant way, he needs to switch to a fast multi purpose language. So in the end he is loosing all the advantages of the high level scientific computing language. A pattern which has evolved out of this dilemma is to prototype in a nice high level language and as soon as the algorithms has been confirmed to work, to rewrite it in a fast low level language. One of the first language promising to solve this dilemma for scientific computing is Julia. It is supposed to be high level and optimized for the work of scientific computing while approaching the speed of statically compiled languages like C. This leads us straight to the field of research and problem that gets solved in this thesis.

1.2 Field of Research and Problem

Figure 1: *different visualizations of* $f(x, y, z) = \sin(\frac{x}{15}) + \sin(\frac{y}{15}) + \sin(\frac{z}{15})$*, visualized with Romeo. From left to right: Isosurface with isovalue=0.76, Isosurface with isovalue=0.37, maximum value projection*

This thesis is about bringing performance and usability together in the realms of scientific computing and 3D visualizations. These two demands are pretty much opposing concepts. One is about bringing tasks into a form of making them best understandable to humans, and the other is about transforming a task to make it fit well to a computer architecture. These two tasks could not be more different. For humans, data and algorithms becomes understandable if they are high level and represented visualy, auditorial or tactile together with immidiate feedback. It is the task of making problems accessible to a human, who has evolved his capabilities in order to survive and find food and not to create complex algorithms. Computers on the other hand love to have their registers filled optimally,

move memory to smaller and faster caches and dislike random access to memory. That is all they care about, whether a human understands this or not.

To close this gaps, compiler have been created. They are translators between human understandable languages to machine instructions. This is just the first step and many more are needed to create an enjoyable user experience. These steps range from introducing graphical user interfaces, novel input devices like the mouse, understandable visualizations and so forth. All these advances have made computers usable for people who do not have an education in computer science. In this thesis the field is scientific computing, which still has quite a lot of barriers for novel users. Scientific computing is usually about implementing mathematical equations, complex algorithms and manipulating and analyzing data. Most research is done in some specialized, high-level scientific computing language. Besides the previously discussed performance problems with this approach, the lack of easily usable, extendable and fast visualization libraries also poses a problem. Most state of the art visualization libraries use C++ at their performance critical core, they are not extendable or they are simple toy libraries, which can not be used for projects with higher performance demands.

This is a problem for several reasons. First, it creates a complexity and performance bottlenecks when interfacing with other languages. The next problem occurs, when the library does not offer the needed functionality and the programmer has to step in and extend the library. Finally, you often do not have easily accessable GUI elements, they come from a different package (possibly written in yet another language) or they are compliacted to use. This makes it hard for the researcher to visualize and interact with his data. It would be desirable, if interactability, speed and extendability would be the default for any visualization library.

Consider the following function $f(x, y, z) = \sin(\frac{x}{15}) + \sin(\frac{y}{15}) + \sin(\frac{z}{15})$, which describes a 3D volume mathematically. This is a simple function, which is already not that easy to interpret. In figure 1, you can see different visualizations of f. If you can interact with this visualization, by moving through the iso values or coloring certain areas, it will make the function more understandable. This deeper understanding is crucial for identifying problems in the underlying math, extending the function, or explaining it to other people. Making problems more understandable like this further opens the gates of scientific computing to novel users.

In summary, the software in this thesis focuses on research which involves writing short scripts, while playing around with some parameters and visualizing the results. An example would be a material researcher, who is investigating different 3D shapes and materials

and their reaction to pressure. The researcher would need to read in the 3D object he wants to analyze, have an easy way to tweak the material parameters and it would be preferable to get instant feedback on how the pressure waves propagate through the object. There are quite a few libraries out there offering this, but none of them is written in a high level language, offers speed, extendability and usability, while being deeply integrated into a scientific programming language.

1.3 Contribution

The main contribution of this thesis is writing Romeo in Julia, which offers the following advantages.

Julia is a high-level language and some effort was also put into creating a concise architecture, so one contribution is, that the development cycles can be very short and the library is easy to extend.

The target group for Julia are researchers, which should be able to write their research completely in Julia. As Julia is fast and the library is also written in Julia, this will enable researchers to stay in the same language for their project. Due to Julia's speed, this makes it easy to create visualisation pipelines in which every routine is as fast as it can be. Also, the researcher can extend the library in the same language he is already working in.

On top of that, the library makes it simple to interact with complex algorithms via widgets and forms a basis for visual debugging. This comes with an ease of use, which would be hard to achieve if the library was not that deeply embedded in Julia.

1.4 Outlook

[short outlook, includes BA structure and some words about the results]

2 Background

In this chapter, a short overview will be given over the current state of the art for visualization inside the field of scientific computing and a short intro to Julia will be given.

2.1 Related Work

2.1.1 The Julia Programming Language

Bringing Julia's ease of use and speed to a dynamic visualization library is the declared goal. So Julia plays a crucial role in this thesis. It is the most important previous work, as much as Julia is the main used technology. This chapter gives a short introduction to the Julia Programming Language.

Julia was published in 2012, which makes it a very new language. It is currently at version 0.3.7 stable and 0.4 pre-release. Following common versioning conventions this means Julia is still in an early release phase with the core features and names suspicable to change. Julia is using the compiler infrastructure LLVM to generate fast assembly code.

Julia is a multi paradigma language for scientific computing. Some of its most important features are multiple dispatch, a dynamic type system, macros, good performance and an interface to C and Python. The focus on scientific computing means, that Julia's standard library is equipped with a lot of functions, data structures and specialized syntax for implementing complex math like linear algebra and statistics. It promises to approach C speed, while being a dynamic language which is easy to use. This is made possible by the compile process which can be described as statically compiled at runtime. Julia uses a garbage collector, taking the task of memory management away from the programmer. There are quite a few things Julia promises to the developer which includes the following items[13]:

- C like performance
- native C interface
- macros like in Lisp
- mathematical notations like Matlab
- good at general purpose programming as Python
- easy for statistics as R.

Another interesting feature of Julia is, that all user defined types are as fast and compact as build in types. This allowes Julia to write all arithmetic types in Julia itself, allowing

anyone to extend and rewrite them without diving into the compiler. This means that you can write your own floating point type, with the same performance and characteristics as the build in floating point type.

Julia claims to take an approach at scientific computing which is more modern than other programming languages. They justify this by pointing out the performance characteristics of Julia and the possibilities that come from this. Julia allows to write even the performance critical core in Julia itself. This means Julia does not need to call out to C and most of the standard library is implemented in Julia. In contrast, Matlab, Python and R need to implement any perfomance criticial code in another language, which has lead to a programming pattern which is known by the name of vectorization. This pattern has evolved, as the vector operations that are build into the language are much faster than self written vector operation. This means, if one needs performance, the code needs to be rewriten to only use functions implemented in some module that relies on another, faster language. A deeper analysis of this problem can be found in the first Julia paper[2].

All in all, this makes Julia a very desirable scientific computing language, which promises to be also great for a visualization library. As part of this thesis, it will be investigated if Julia's claims have been achieved.

2.1.2 LLVM

LLVM is an compiler infrastructure, which has front ends for different languages and compiles to different platforms like x86, ARM, Open Compute Language (OpenCL) and Compute Unified Device Architecture (CUDA). A language designer must create an Abstract Syntax Tree (AST) which than LLVM can convert into LLVM Intermediate Representation (IR). This IR can than be heavily optimized. Every language that can be converted to LLVM IR can be combined at this level, going through the same optimizations in the end. This is yields superiour language interfacing, as inlining and other optimizations can be done over the boundary of one language.

LLVM's concept is effective, as you can accumulate state of the art optimizations in one place, making them accessible to many languages. Because of the many backends, the language can run on many architectures. While Julia does not support them all, it will hopefully be possible in the future. LLVM is also used by Clang, the C/C++ front end for LLVM rivaling gcc and it is used by Apple's programming language Swift. This makes LLVM a solid basis for a programming language, as these are highly successful projects guaranteeing LLVM further prospering.

2.1.3 IJulia

IJulia is the Julia language back-end for IPython. IPython is a software stack, which was created to allow for interactive computing in Python. It offers an interactive shell to execute python scripts, GUI toolkits, tab completion and rich media visualizations. It comes with a web based notebook, which enables you to write formated documentations together with data, inlined plots and executable program snippets. You can also formulate mathematical formulas in latex, which will get rendered and inlined nicely into an IJulia Notebook. See figure 16 for an example.

IJulia has some similar goals compared to Romeo, but it has a different focus. The notebook is completely web based, concentrates on 2D visualizations and interactivity is mostly limited to the programming and not the graphics. 3D graphics are possible via Three.js, which is a powerful 3D visualization library based on WebGL. The current integration is just prototypical and limited to simple 3D meshes up to now.

2.1.4 Matlab

Matrix Laboratory (Matlab) is a numerical computing environment that comes with its own programming language. It was created in 1984 by Cleve Moler. He designed it to leverage the effort of accessing LINPACK and EISPACK for his students. Since then it grew to be a widely used tool for scientific computing in all areas, ranging from teaching to actual engineering uses in companies. It offers a broad range of functionality, including matrix manipulation, plotting of functions and data, creation of user interfaces and interfacing with a range of languages like C/C++, Java Fortran and Python. Matlab has made itself quite the name with having print ready visualization tools deeply integrated into the standard library.

Matlab itself is written in C, C++, Java and MATLAB. It's proprietary software with a pricing of around 2000 €[21], which can be extended via free, open source and proprietary modules like Simulink.

Romeo intends to lay out the ground work to provide a similar deep integration of visualizations in Julia. It is quite far away in terms of functionality, but it builds upon a more modern architecture. Romeo is using modern OpenGL and Julia intends to solve one of matlabs biggest problems, namely the need for vectorization. Overall, the biggest difference is that Romeo and Julia are open source, making them much more accessible and easier to extend than Matlab.

2.1.5 Mayavi and VTK

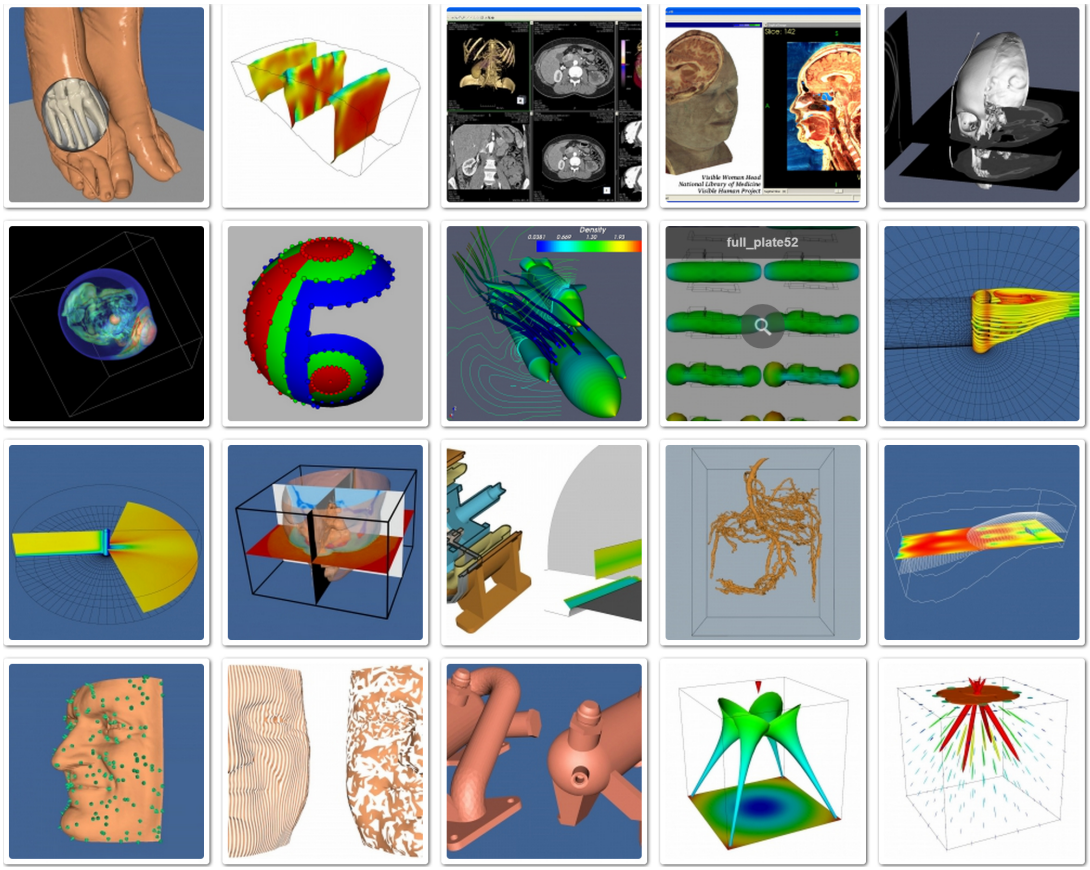

Figure 2: *Different visualizations done with VTK.*

Mayavi is probably one of the biggest, open source library for interactive 3D visualizations. It is written 99.9% in python, but relies on Visualization Toolkit (VTK) for rendering. VTK is one of the most advanced scientific visualization library, with a huge amount of visualization types. In figure 2 you can see some of the visualization taken from the VTK gallery[14].

Mayavi shares some of its goals with Romeo, namely[5]

- An (optional) rich user interface with dialogs to interact with all data and objects in the visualization.
- A simple and clean scripting interface in Python, including one-liners, or an object-oriented programming interface. Mayavi integrates seamlessly with numpy and scipy for 3D plotting and can even be used in IPython interactively, similarly to Matplotlib.
- The power of the VTK toolkit, harnessed through these interfaces, without forcing you to learn it.

Obviously, the python part is not a shared goal, but creating an interactive 3D visualization library deeply embedded into a language is. Mayavi together with VTK is a very big project and in this sense not really comparable to Romeo. It amounts to a total of 3.642.105 lines of code written in 29 languages. The statistics can be found in table 8 and 10. The biggest difference is, that Romeo is implemented in a scientific programming language, while Mayavis core uses VTK which is mainly implemented in C++. This has two big implications. Firstly, if the language does not have native C++ compatible data types and an overhead less C++ interface, shipping a large stream of data to VTK becomes slow. Secondly, one must know C++ to extend VTK. This makes it difficult to create customized visualizations.

In contrast, Romeo is implemented in one language, making these tasks very simple and efficient.

2.1.6 Vispy

Vispy is yet another interactive 3D visualization library. It is from the goals and development status the closest to Romeo. These include[6]:

- High-quality interactive scientific plots with millions of points.
- Direct visualization of real-time data.
- Fast interactive visualization of 3D models (meshes, volume rendering).
- OpenGL visualization demos.
- Scientific GUIs with fast, scalable visualization widgets (Qt or IPython notebook with WebGL).

It is a fairly new library, promising to use modern OpenGL and state of the art performance. This is very similar to Romeo's goals, with the only difference being that Romeo is implented in Julia while Vispy is implemented in Python. So the biggest differenciation between Romeo and Vispy will be the performance and the concrete feature set.

3 Requirements

All building blocks in this thesis are developed with the purpose in mind to give the user the possibility to visualize and interact with complex 2D and 3D data, while being able to easily extend the library. To enable this kind of functionality, a lot of parts of the infrastructure need to work seamlessly together. Certain design choices had to be made to guarantee this. As speed is the most constraining factor, this chapter will start by introducing the design choices that had to be made in order to achieve state of the art speed.

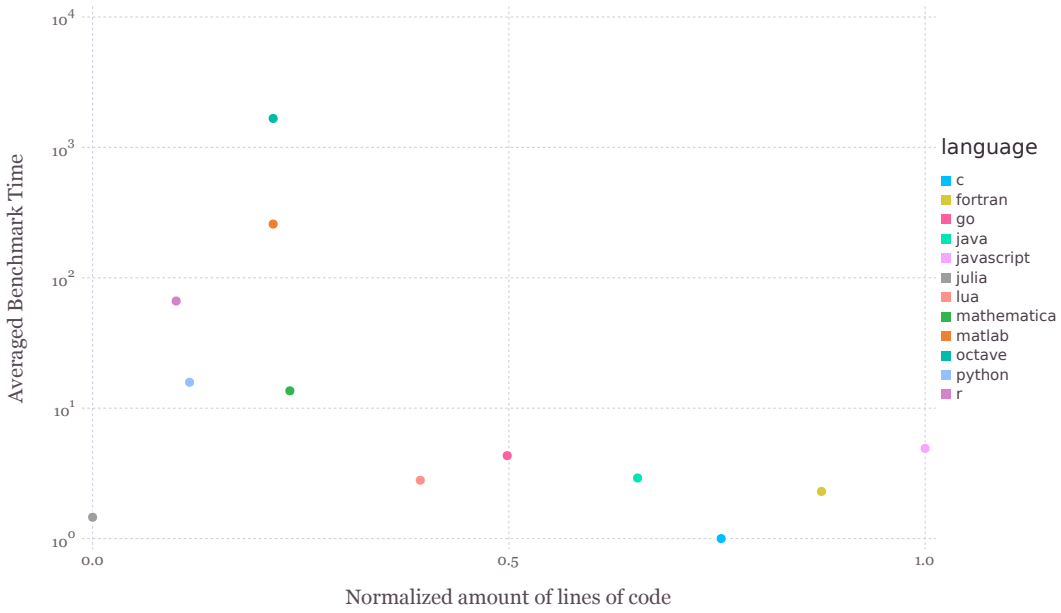

Figure 3: *Languages speed relative to C (averaged benchmark results), plotted against the length of the needed code (Source in Appendix).*

3.0.7 Speed

Speed is mainly a usability factor. It is a factor, that can make a software unusable, or render it unproductive. Because of this, speed has taken a high priority in this thesis. As general coding productivity is also a concern, this thesis is set on using a high level language. Historically, these two demands can not be satisfied at the same time. How to achieve state of the art speed with a high level language is an ongoing research and basically the holy grail of language design. Julia promises to do exactly this, which is illustrated in figure 4. Code length is an ambiguous measure for conciseness, but if the code is similarly refactored it is a good indicator of how many lines of code are needed to achieve the same goal. From this figure we can conclude, that Julia at least comes close to its promises, which is why it has been chosen as the programming language.

To get high performant 3D graphics rendering, there are on the first sight a lot of options. If you start to take the previous demands into account, the options shrink down considerably. The visualization library should be implemented in one high level language, which can be used for scientific computing and has state of the art speed. At this point, there are close to zero libraries left. As you can see in figure 4, Matlab, Python and R disqualify, as they are too slow. JavaScript, Java, Go and Lua are missing a scientific background and the others are too low level for the described goals. This leaves only Julia, but in Julia there were no 3D libraries available, which means that one has to start from scratch. There are only a couple of GPU accelerated low-level libraries available, namely Khrono's OpenGL, Microsoft's DirectX, Apple's Metal and AMD's Mantel, which are offering basically the same functionality. As only OpenGL is truly cross-platform, this leaves OpenGL as an option. So for the purpose of high speed visualizations, OpenGL was wrapped with a high-level interface written in Julia. This leaves us with one binary dependency not written in Julia, namely the video driver, which implements OpenGL.

Measurement of success is pretty straight forward, but the devil is in the detail. It is easy to benchmark the code, but quite difficult to find a baseline, as one either has to implement the whole software with an alternative technologies, or one has to find similar software. This thesis will follow a hybrid strategy, comparing some simple implementations with different technologies and choose some rivaling state of the art libraries as a baseline.

3.0.8 Extensibility

Extensibility is an important factor, which can decide, if a library is fit for scientific computing or not. It is not only that, but also a great factor determining growth of a software, as the more extensible the software is, the higher is the probability that someone else contributes to it. In order to write extensible software, we first have to clarify what extensibility is. Extensible foremost needs that the code is accessible. There are different levels of accessibility. The lowest level is closed source, where people purposely make the code inaccessible. While this is obvious, it is just a special case of not understanding the underlying language. Just shipping binaries without open sourcing the code, means that the source is only accessible in a language which is extremely hard to understand, namely the machine code of the binary. So another example for inaccessibility is to write in a language that is difficult to understand. Other barriers are obfuscated language constructs, missing documentations and cryptic highly optimized code. Further more the design of the library in the whole is an important factor for extensibility. It is not only important, that all parts are understandable, but also, that every independent unit in the code solves only one problem. This guarantees that one can quickly exchange it, or use it somewhere else where the same problem needs to be solved. If this is guaranteed,

re-usability in different contexts becomes much simpler. This allows for a broader user base, which in turn results in higher contributions and bug reports. Short concise code is also important, as it will take considerably less time to rewrite something, as the amount of code that has to be touched is shorter and less time is spend on understanding and rewriting the code.

So the code written for this thesis will be open source, modular, written in a high level language and concise.

This is pretty difficult to measure as these are either binary choices, which are followed or not, or higher level concepts like writing concise code, which can be a matter of taste. To get an idea of the effectiveness of the strategy, usage patterns and feedback from Github will be analyzed.

3.0.9 Event System

The event system is a crucial part of the library, as the proclaimed goal is to visualize dynamic, animated data. This means, there are hard demands for usability and speed on the event system. The chosen event system has an immediate influence on how to handle animations. This lead to the design choice of using signals. Signals are a very good abstraction for values that change over time. If well implemented, it makes it natural to reason about time, without the need of managing unrelated structures and callback code.

3.0.10 Interfaces

Working with a computer means working with interfaces to a computer, which in the end simply juggles around with zeros and ones. There is a huge hierarchy of abstractions involved, to make this process of binary juggling manageable to the human. We already dealt with the lowest relevant abstraction: the choice of programming language, which forms our first interface to the computer. The next level of abstraction is the general architecture of the modules, which has been discussed previously. This chapter is about the API design choices that have been made.

The first API is the OpenGL layer. The philosophy is to make the wrapper for native libraries as thin and reusable as possible and an one to one mapping of the underlying library. This guarantees re-usability for others, as they might be used to work only with the low-level library or they might disagree with some higher-level abstraction and prefer to write their own.

Over this sits an abstraction layer needed to simplify the work with OpenGL. With this abstraction, the actual visualization library is implemented.

API! (**API!**)s for visualization libraries are very difficult to realize, as there are endless ways of visualizing the same data. The design choice here was to use Julia's rich type system to better describe the data. Julia makes this possible, as you can create different types for the same data, without loosing performance. So you can have a unit like meters represented as a native floating point type and have the visualization specialize to this. Like this you can have a single function e.g. *visualize*, that does create a default visualization for different data types. Instead of manually passing additional information to the visualization function, it is coded in the type itself. Together with the event system which consists of signals, it is possible to edit and visualize rich data over a simple interface, which is perfect for visual debugging, as it is always the same function call applied to the data and no further user interaction is needed. It is also easy to extend, as the user just has to overload the function with a custom style and optional key word arguments. Finally, there are also graphical user interfaces developed for this thesis. As also optimizing them is out of the scope of this thesis, they are kept very simple. The measurement of success is again relatively difficult to do. (I need to think this over)

4 Used Technologies

4.1 The Julia Programming Language

The basic introduction of Julia has already been given in the Background chapter. This chapter is focused on how to write programs with Julia. Most influential language construct are its hierarchical type system and multiple dispatch. Multiple dispatch is in its core function overloading at runtime. To better understand multiple dispatch, one has to be familiar with Julia's type system. The type system builds upon four basic components. Composite types, which are comparable to C-Structs, parametric composite types, bits types, abstract and parametric abstract types. While the first three are all concrete types, abstract types can not be instanciated but are used to build a type hierarchy. Every concrete type can only inherit from one abstract type, while abstract types can also inherit from abstract types. Bit types are just immutable, stack allocated memory chunks, usable for implementing numbers. You can build type hierarchies like this:

```
1 abstract Number
2 abstract FloatingPoint{Size} <: Number # inherit from Number
3 bitstype 32 Float32 <: FloatingPoint{32} # inherit from a parametric
      abstract type
4 type Complex{T} <: Number
5     real::T
6     img::T
7 end
```

With this type hierarchy you can overload functions with abstract, concrete or untyped arguments.

```
1 foo{T}(y::Complex{T}, y::Float32) = println("some number: ", x, " some
    complex Number: ", y) # shorthand function definition
2 function foo(x)
3     println("overloading foo with a new unspecific signature")
4 end
```

What will happen at runtime is, that Julia compiles a method specialized on the arguments which results in overloading the function with the concretely typed arguments. To illustrate this let us look at the example. Initially, *foo* will be overloaded with two methods. Now, if you call *foo* with one *Float32* argument, a new method will be added at run time specialized to *Float32*. Like this, if the function does not access non constant global values, all types inside the function will be known at call time. This allows Julia to statically compile the function body, while propagating the type information down the call tree.

With multiple dispatch, Julia is a functional oriented language. But there are also ways to give Julia a more object oriented feel. Functions are first-class, so they are easy to pass around. They can be bound to variables and can then be called like normal functions via the variable name. This implies that functions can also be bound to objects. There is no self reference available in Julia, so the object still needs to be passed to function via the function arguments

Another of the most crucial features is the very simple, overhead less C-Interface. Thanks to the binary compatibility of LLVM's emitted assembly, a C function call to a shared library inside Julia has the same overhead as it would be from inside C[15]. This is perfect for integrating low-level libraries like OpenGL and OpenCL.

Julia's performance is crucial for this thesis. If Julia does not perform close to C it would weaken the whole argument of writing the visualization library in Julia. It is a very tedious task to write representative benchmarks for a programming language. The only way out is to rely on a multitude of sources and try to find analytical arguments. In this thesis, Julia's own benchmark suite will be used in addition to some real world benchmarks. In addition, the general compiler structure of Julia will be analyzed to find indicators for the limits of Julias performance.

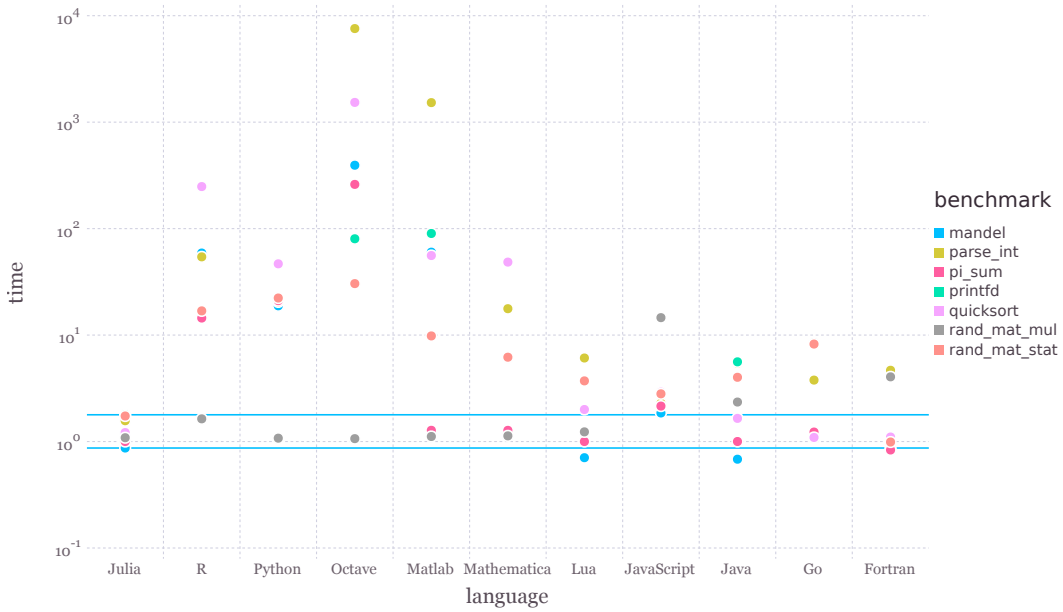

Figure 4: *Julia's performance compared to other languages, taken from Julia's micro bench suite [16]. Smaller is better, C performance = 1.0.*

In the first benchmark from figure 4, we can see that Julia stays well within the range of C Speed. Actually, it even comes second to C-speed with no other language being that close. This is a very promising first look at Julia, but it should be noted, that these benchmarks are written by the Julia core team. So it is not guaranteed, that there is no bias favoring Julia in these benchmarks. There is another benchmark comparing C++, Julia and F#, which was created by Palladium Consulting which should not have any interest in favoring one of the languages. They compare the performance of C++, Julia and F# for an IBM/370 floating point to IEEE floating point conversion algorithm. This is part of a blog series[9] written by Palladium Consulting. F# comes out last with 748.275 ms, than Julia with 483.769 ms and finally C++ with 463.474 ms. At the citation time, the Author had updated the C++ version to achieve 388.668 ms. It looks like the author was only working on making the C++ version faster, so it can not be said that the other versions could not have been made faster too.

The last Julia benchmark is more real world oriented. It is comparing Finite Element solver, which is an often used algorithm in material research and therefore represents a relevant use case for Julia.

These are remarkable results, considering that the author states it was not a big effort to achieve this. After all, the other libraries are established FEM solvers written in C++, which should not be easy to compete with.

N	Julia	FEniCS(Python + C++)	FreeFem++(C++)
121	0.99	0.67	0.01
2601	1.07	0.76	0.05
10201	1.37	1.00	0.23
40401	2.63	2.09	1.05
123201	6.29	5.88	4.03
251001	12.28	12.16	9.09

Table 1: *Performance of a FEM solver written in Julia compared to some existing libraries. [24]*

This list could go on, but it is more constructive to find out Julia's limits analytically. As already mentioned, Julia is statically compiled at runtime. This means, as long as all types can be inferred at runtime, Julia will have in the most cases identical performance to C++. The biggest remaining difference in this case is the garbage collection. Julia 0.3 has a mark and sweep garbage collector, while Julia 0.4 has an incremental garbage collector. As seen in the benchmarks, it does not necessarily introduce big slowdowns. But there are issues, where garbage collection introduces a significant slow down[23]. Analysing this further is not in the scope of this thesis, though. But it can be said that Julia's garbage collector is very young and only the future will show how big the actual differences will be.

Another big difference is the difference in between different compiler technologies. LLVM's biggest rival is gcc. If C++ code that is compiled with gcc is much faster than the same code compiled with LLVM, the gcc version will also be faster as a comparable Julia program. In order to investigate the impact of this, one last benchmark will be analyzed. This is a summary of a series of articles posted on Phoronix, which benchmarked gcc 4.92 against LLVM 3.5 and LLVM 3.5 against LLVM 3.6:

Statistic	gcc vs LLVM 3.5	LLVM 3.5 vs LLVM 3.6
mean	0.99	0.99
median	0.97	1.00
maximum	1.48	1.10
minimum	0.39	0.88

Table 2: *Summary of the Phoronix benchmark. Unit is speedup of LLVM, bigger is better. [1][17][18]*

The full tables can be found in the appendix under the table 14 and 16. The results suggest, that LLVM is well in the range of gcc, even though that there can be big differences between the two. These are promising results, especially if you consider that LLVM is much younger than gcc. With big companies like Apple, Google[28] and Microsoft[19] being invested in LLVM, it is to be expected that LLVM will stay competitive, which

means Julia should in theory stay competitive as well.

4.2 OpenGL

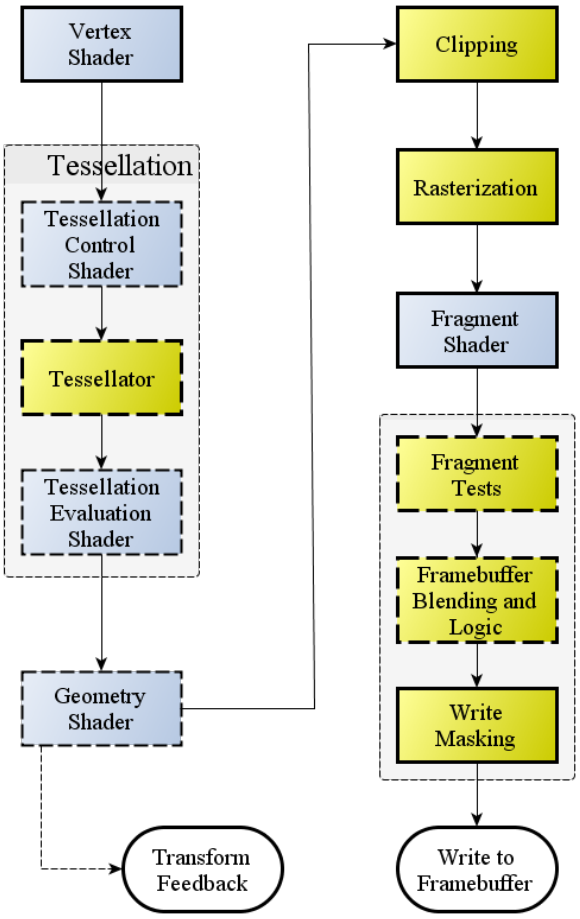

Figure 5: *Diagram of the Rendering Pipeline. The blue boxes are programmable shader stages. Arrows show the flow of data[26]*

OpenGL is a low-level graphics API implemented by the video card vendor via the video driver. As such it does not offer much abstraction over the actual Graphics Processing Unit (GPU), but instead offers high flexibility and performance. OpenGL 1.0 was released in 1992 and the current version is 4.5. A critical element when developing OpenGL applications is, that not all video drivers implement the newest OpenGL standards.

As a result, one has to decide which OpenGL version to program against, trading between modernity and platform support. For Romeo, it was decided to support OpenGL 3.3 as the lowest bound, as it is sufficiently available, while still having most of the modern features. All features that help to call less OpenGL functions can be considered as modern, as they take away the load from the CPU. The modern features used in this thesis include instance rendering, vertex arrays and OpenGL Shading Language (GLSL) shader.

In figure 5 you can see the basic architecture of an OpenGL program pipeline. As the description states, the blue boxes are programmable shaders, while the dotted boxes are optional parts of the pipeline. The yellow boxes describe stages which are not directly accessible. They are part of the global OpenGL state, which can be set via OpenGL commands.

So in order to have a functioning OpenGL rendering pipeline one needs to write at least a vertex shader and a fragment shader. All shaders are compiled and linked into a program object, which then can be executed on the GPU. Shaders are written in a C dialect specialized for vector operations. You feed shaders with data via buffers, textures and uniforms. Buffers are 1D arrays, textures 1D/2D/3D arrays with both having their own memory, while uniforms live in the program object.

The different shaders are usually used to apply geometric, perspective transforms and calculating the light. In newer APIs general compute operations are available, making it possible to create more flexible shader stages. Finally, the fragment shader rasterizes the data to the render targets in the frame-buffer. The frame-buffer can then be displayed on the monitor. Frame-buffers can contain multiple render targets, which are buffers that the fragment shader can write to. The write operation is heavily restricted. The fragment shader can only write to the location calculated by the vertex shader and simultaneously reads from the frame-buffer are not possible. This restriction exists to allow for the massive parallel execution model that OpenGL uses to speed up rendering times.

The usual set of render targets includes a depth channel, stencil buffer and of course the color buffer. The depth channel is usually used to discard all fragments that are behind another fragment, while the stencil buffer can be used to discard arbitrary fragments. Custom render targets can be created in newer OpenGL versions which can be directly addressed via the fragment shader. Here is a simple minimal example for a program rendering some vertex data with a flat color to the screen.

```
1  //Vertex Shader
2  in vec3 vertex; // vertex fed into the shader via a buffer
3  uniform mat4 projection; // Projection matrix
4  uniform mat4 view; // View matrix, setting rotation and translation of
       the camera
5  void main()
6  {
7      gl_position = projection*view*vec4(vertex, 1); // apply
           transformations to vertex and output to fragment shader
8  }
9  //Fragment shader
10 out vec4 framebuffer_color; //color render target, which will get
```

```
        written into the display framebuffer
11  void main()
12  {
13      framebuffer_color = vec4(1,0,0,1); // write a red pixel at
            gl_position from the vertex shader.
14  }
```

All visualization code is written in OpenGL shaders, which are compiled and executed via GLAbstraction.

4.3 Reactive

Reactive[10] is a functional event system designed for event driven programming. It implements Elm's[3] signal based event system in Julia. Signals can be transformed via arbitrary functions which in turn create a new signal. This simple principle leads too a surprisingly simple yet effective way of programming event based applications.

```
1  a = Input(40)          # an integer signal.
2  b = Input(2)           # an integer signal.
3  c = lift(+, a,b)       # creates a new signal with the callback plus. Equal
       to c = a+b
4  lift(println, c)       # executes println, every time that c is updated.
5  push!(a, 20)           # updates a, resulting in c being 22
6  #prints: 22
7  push!(b, 5)            # updates a, resulting in c being 22
8  #prints: 25
```

Lifting a signal creates a callback, which gets called whenever the signal changes. There are more operations than lifting, like folding, merging, filtering and so on. With this, one can build up a complex tree of operators which will get applied to the origin signal. For the concrete case of Reactive, every signal carries around a list of children and parents. Each signal has a rank, in order to build up a sorted heap with these information. So every time a signal is updated, the heap can be traversed and the functions get applied in the right order, updating all the values of the children. Reactive is used in all parts of the library. It builds the basis for the camera code, the widgets and any value that needs to be animated is realized via a signal.

4.4 GLFW

GLFW[8] is a cross platform OpenGL context and window creation library written in C. GLFW allows to register callbacks for a multitude of events like keyboard, mouse and window events. This, together with a wrapper library for Julia makes GLFW perfect

for doing the window creation. In addition, GLFW exposes low level features like the operating systems context handle. This can be used for creating advanced contexts that share memory with another context. Romeo does not use this feature yet, but it makes GLFW a future proof choice.

5 Implementation

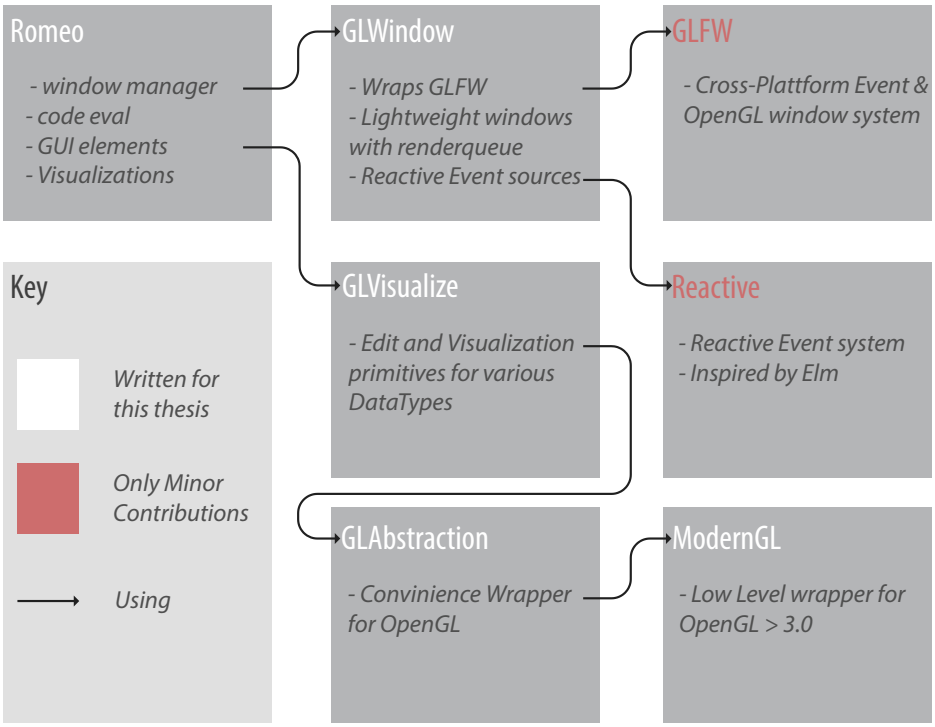

Figure 6: *Main modules used in Romeo and their relation (simplified).*

This chapter is about the implementation of Romeo. The Romeo package itself is small and just defines the high-level functionality of the editor. This includes window layout and connecting all the different event sources to create the wanted behavior. To do this, Romeo relies on a multitude of packages, which step for step abstract away the underlying low-level code that is used to do the window creation and rendering. As you can see in figure 6, Romeo uses GLVisualize for creating GUI elements and the visualisations. The code evaluation is done via Julia build-in functions. Windows are managed with GLWindow. GLWindow creates an OpenGL window with the help of GLFW and converts all window events into Reactive's signals. It also offers a very simple render queue, for rendering graphics attached to a window. Reactive signals are not only used as the event sources, but are also the main abstraction for time varying values in GLVisualize. GLVisualize is the main package offering the rendering functionality and the editor widgets like text

fields and sliders.

For rendering GLVisualize relies on GLAbstraction, which defines a high-level interface for ModernGL. ModernGL does the OpenGL function loading and exposes all the function and Enums definitions from OpenGL with version higher than 3.0. Like already pointed out in the requirements, special care has been taken to make all modules self sufficient. Every single package can be used for other applications, which allows for higher flexibility and a broader user base.

5.1 Event System

The event system was challenging to integrate for several reasons. First of all Reactive is a functional event system, while OpenGL relies heavily on global states, which are two perpendicular concepts. Also, it does not allow to rearrange the event tree. In other words, you can not create sub trees in advance and then fuse them together at run time.

5.2 ModernGL

OpenGL is implemented by the video card vendor and is shipped via the video driver, which comes in the form of a C library. The challenge is, to load the function pointers system and vendor independent. Also one further complication is, that depending on the platform, function pointer are only available after an OpenGL context was created and may only be valid for this context. [22] This problem is solved, by initializing a function pointer cache with null and as soon as the function is called the first time the real pointer gets loaded.

The OpenGL function loader from ModernGL has undergone some changes over the time. Starting with a very simple solution, there have been pull requests to include better methods for the function loading. The current approach in ModernGL master was not written by myself, but by the Github user aaalexandrov. Before aaalexandrov's approach, the fastest approach would have used a pretty new Julia feature, named staged functions. It should in principle yield the best performance as it compiles a specialized version of the function when it gets called for the first time. This is perfect for OpenGL function loading, as the pointer to the function can only be queried after an OpenGL context has been created. When the staged function gets called the pointer can be queried and gets inlined into the just in time compiled function.

Staged functions only work with the newest Julia build, which is why aaalexandrov's approach was used in the end.

5.3 GLAbstraction

GLAbstraction is the abstraction layer over ModernGL. It wraps OpenGL primitives like Buffers and Textures in Julia objects and hooks them up to Julia's garbage collector. Additionally, it implements convenient functions to load shader code and it makes it easy to feed the shader with the correct data types. Besides supplying an abstraction layer over OpenGL, it also offers the linear algebra needed for the various 3D transformation and camera code. Building up on that, it defines a signal based perspective and orthographic camera type.

5.4 GLWindow

GLWindow is a lightweight wrapper around GLFW. It mainly offers a screen type, which contains signals for all the different GLFW events. It also offers a hierarchical structure for nesting screens. All the screen areas are signals, which makes it easy to change the screen area. This makes it simple to implement windows that react to changing the size of the windows or resized objects.

5.5 GLVisualize

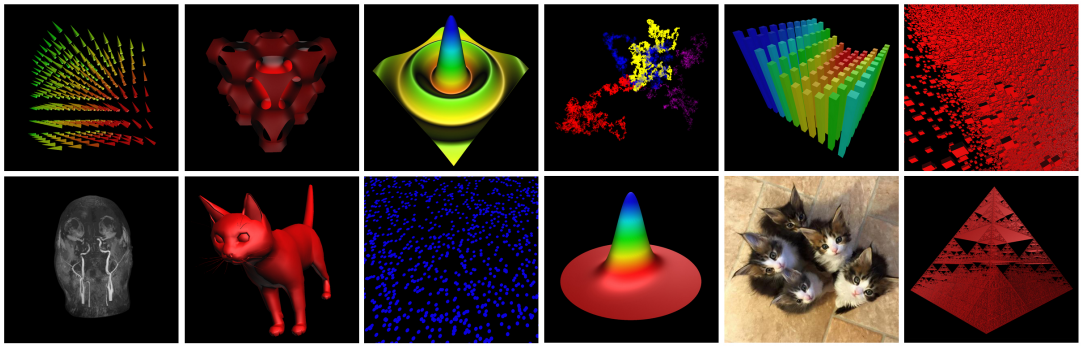

Figure 7: *Different visualizations rendered with GLVisualize.*

GLVisualize implements the main functionality of this library. It structure is quite simple. It relies as much as it can on common Julia data types and creates specialized visualizations for them via dispatch. So instead of offering differently named functions for different visualizations, there is just one function with lots of methods for different types. This has two advantages. First, it makes it very easy to use for visual debugging, as any value can be displayed immediately without any user interaction. Secondly, the user does not have to remember or lookup the function name, as long as there is a default visualization for the type he is working with. The next design goal was to make this fit for dynamic data, which resulted in relying on as little transformation of the data as possible and directly

transferring it to the GPU. Depending on the complexity of the visualization, this means the visualization can be updated with as little overhead as possible.

The interface to create visualizations is very simple and only consists of three functions:

```
1 Dict{Symbol, Any}      = visualization_defaults(data::Union(Signal{T}, T)
      , style::Style) # returns a dictionary of parameters
2 RenderObject       = visualize(data::Union(Signal{T}, T), style=Style{:
      default}; parameters...) # returns an object which can be directly
      rendered
3 RenderObject, Signal  = edit(data::Union(Signal{T}, T), style=Style{:
      default}; parameters...) # returns an RenderObject and signal which
      outputs the changed values
```

With this simple interface, the following data can be visualized:

- Text (Vector of Glyphs)
- Height fields with different primitives (Matrix of height values)
- 3D bar plots (Matrix of height values)
- Images (Matrix of color values)
- Videos (Vector of Images)
- Volumes (3D Array of intensities)
- Particles (Vector of Points)
- Vector Fields (3D array of directional Vectors)
- Colors (Single Color values)

All of these can be integrated into the same scene and it is possible to change their parameters interactively. These interactions can be purely programmatically, or via the widgets from the edit function. It calls the visualize function to render the data type and then registers appropriate events to update the data. Take a look at the text edit function. It first uploads the text to video memory and sets up the functionality to visualize it, and than updates the text data on the GPU according to the cursor position and keyboard input.

Up to now, there is an edit function available for strings, colors, numbers, vectors and matrices.

5.5.1 Mesh primitive Rendering

The rendering of meshes in OpenGL is pretty straight forward with a normal vertex and fragment stage. Vertex, Normal and UV data is supplied via Vertex Arrays, perspective transformations via uniform matrices. In the fragment shader, a Blinn-Phong lighting model is applied.

5.5.2 Particle Rendering

Most of the visualizations in GLVisualize are realized via instancing a mesh primitive. So the bar plot is nothing else than a cube placed in a grid, with scaling informations that get applied to every individual cube. The surface plot is a quad or any other 2D mesh spaced across a grid, while the vertexes are projected onto a height field. The vector field is a mesh placed regularly inside a cube, while the rotations from the vector field gets applied to this mesh. Even text rendering functions in the same way. The difference is just, that the particle not only holds position information, but also indexes to a texture atlas in which renderings of the glyphs are cached. So when rendering the text particles, the exact scale and image of the glyph is queried, which will then be used to render a quad with the image of the particular glyph to the screen. The texture atlas approach was chosen, because rendering a high quality vector graphic is very time consuming, especially if the description of the font is only available as a Bézier spline. A more detailed description can be found in 5.5.3. All particles are rendered via OpenGL's instanced rendering API, which allows to render millions of particles with only one draw call and very little memory usage, as the geometry of the particle just needs to be uploaded one time. For every individual particle additional information like color, position, scale and so forth can be queried from within the fragment or vertex shader stage. This additional information can be stored in uniform buffers, uniform arrays or textures. Textures have been chosen for this thesis, as they offer the greatest support among devices and are easy to use. In the future, other approaches can be implemented, gaining more performance or flexibility. The texture approach has the disadvantage that 1D textures only offer a maximum sizes between 1024 and 8192 elements, so for greater amounts of particles a 1D vector has to be transformed to a 2D texture.

5.5.3 Vector Graphics Rendering

```
From a speech of Demosthenes in the 4th century BC:

Οὐχὶ ταὐτὰ παρίσταταί μοι γιγνώσκειν, ὦ ἄνδρες ᾿Αθηναῖοι,
ὅταν τ᾿ εἰς τὰ πράγματα ἀποβλέψω καὶ ὅταν πρὸς τοὺς
λόγους οὓς ἀκούω· τοὺς μὲν γὰρ λόγους περὶ τοῦ
τιμωρήσασθαι Φίλιππον ὁρῶ γιγνομένους, τὰ δὲ πράγματ᾿
```

Figure 8: *GLVisualize's fast UTF8 rendering with the help of Cairo, a texture atlas and 2D particles.*

Vector graphics are difficult to render, as they are not well fit for the GPU. Long stretched, curved and thin lines introduce several problems for the GPU[20]. Another problem is that splines used in vector graphics are usually supplied as Bézier curves, which are very demanding to rasterize. Besides from that, anti-aliasing of thin lines introduces some problems as well. With post processing anti-aliasing techniques, lines which are thinner than one pixel will introduce artifacts, as OpenGL primitives smaller than a pixel will get discarded by the OpenGL pipeline. So additional care needs to be taken in order to assure, that primitives are always thicker than one pixel, or multi-sampling techniques have to be used. This is only a very short summary of the problems, which is only given to illustrate that this is not a problem that can be solved in the scope of this thesis. Instead, Cairo is used for rendering fonts and other vector graphics. As Cairo is relatively slow, these renderings get cached in a texture atlas from which they can get queried and rendered to the screen in large numbers. This results in high quality and fast renderings. This comes at the cost of higher space requirements and resolution dependence. So when zooming into the vector graphics, either new versions have to be rendered with Cairo, or one gets pixelated results. In the future, distance fields can be used to reduce the resolution dependence[11].

5.5.4 Volume Rendering

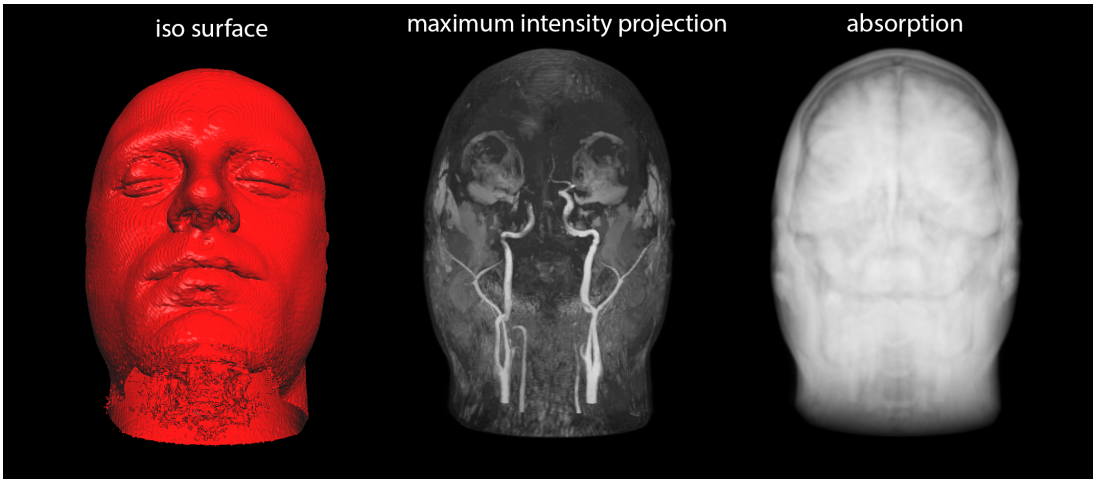

Figure 9: *Different visualizations rendered with GLVisualize. As can be seen, every methods misses critical features of the volume.*

Volumes in 3D computer graphics are usually represented with voxels, which can be understood as 3D pixels. They represent values on a 3D grid. The name stems from the marriage of volume and pixel. There are sparse and dense storage options for voxels and different informations can be represented, like speed, rotation, color, intensities and so forth. Volume renderings are no trivial tasks. Not only from a computational point of view, but it is also difficult to make a visualization which lets you peek inside a volume without discarding important informations. This is illustrated in figure 9. This is why different forms of visualizations are needed to get a good representation of a volume. GLVisualize is able to render iso-surfaces, maximum intensity projections and an absorption based visualization form. On top of this, the particle systems can be used to give further insights into the volume by rendering particles for each voxel. This is especially helpful for sparse volumes, as these are not yet supported natively by GLVisualize.

For the Volume rendering two techniques are being used. Marching tetrahedra and ray marching. Marching tetrahedra is an algorithm that can be used to extract a mesh representation of an iso surface from a volume. Ray marching is the process of shooting rays from the camera origin through the object. With every step inside the volume, values are sampled and depending on the combination of these values, different visualization forms can be realized. When you stop at a certain value, iso-surfaces can be rendered. If only the maximum is kept, it will yield a maximum intensity projection. When all values are combined via an absorption function, the volume will be displayed as if it is made of some translucent material like smoke.

The ray marching rendering method was explicitly implemented for this thesis on the

GPU. For the marching tetrahedra algorithm there was already a Julia implementation available in the package Meshes.jl[25]. The resulting mesh can be displayed with GLVisualizes mesh rendering capabilities. Both techniques have their advantages and disadvantages. While marching tetrahedra is relatively slow, it can be used to generate a mesh which is very fast to display. Ray casting on the other hands allows for a wide range of visualization forms and is rather fast as it runs on the GPU. The downside is, that the calculation can not be cached camera position independent.

5.6 Scene Graph

The scene graph is in the case of Romeo not a specialized data structure, but rather just a list of objects, which can be directly rendered with OpenGL. Functionality from Reactive, GLAbstraction, GLVisualize and GLWindow are involved in this. GLvisualize creates a renderable object with GLAbstraction, which will get pushed into a render queue in the screen from GLWindow. Everything that moves is handled via signals from Reactive. This is an extremely simple form of a scene graph, which does not allow to perform any optimizations. Optimizations usually include sorting in order to reduce OpenGL state changes and culling of invisible objects. As GLVisualize produces OpenGL code which relies only on very few calls to OpenGL, the first optimization is currently not as important. But culling can make a large difference for big 3D scenes, as usually only a fraction needs to be rendered. As this is a more involved process, which would preferably be done completely on the GPU, this was not in the scope of this thesis.

5.7 3D Picking

3D picking is the process of inferring, what object belongs to the pixel on the screen. It forms the basis for any mouse interaction with objects displayed on the screen. To make this as simple and fast as possible, a fairly simple approach has been chosen. The approach is called color picking. For color picking, two frame-buffer render targets need to be created. One for the color channel and one to represent the object id plus an additional number to store contextual information. The additional number is usually used for the instance index, which can be used to e.g. infer what text glyph is selected. When rendering, the fragment shader does not only write the color into the render target, but if the color is opaque also the current object id. This way transparency aware 3D picking can be achieved without an extra processing step. The advantage of this methods is that one does not need an extra pass over the geometry. The disadvantage is higher space requirements and that OpenGL's native anti-aliasing does not work well together with the additional render target. Also, the OpenGL pipeline has to be flushed in order to read from the frame-buffer. The anti-aliasing problem can be solved by implementing an

extra anti-aliasing post processing step. This is the way to go, as OpenGL's native anti-aliasing is not very efficient. But as there was no time to do this, it means GLVisualize does not offer any anti-aliasing for mesh rendering in the current state.

5.8 Romeo

So far Romeo just consists of one file with 500 lines of code. It just defines some simple text field, a search field, and a visualize and edit window. The texts gets evaluated as Julia code as soon as it changes. Like this, the text field acts like a very simple Read Eval Print Loop (REPL). Via the search field, you can execute simple Julia statements and the results will be displayed in the visualize window, while all parameters can be edited via the edit window. This means, if you type in a simple variable, the variable will be visualized. But you can also search and transform a variable via simple Julia terms.

6 Results and Discussion

6.1 Performance Analysis

This chapter supplies some benchmarks, to analyze how close this thesis comes to achieve the wanted performance, which should be on eye level with C.

If not stated otherwise, benchmarks are written for this thesis and executed on an Intel Core i5-4200U with an HD 4400 graphic and 8GB of RAM. Julia 0.4 has been used, C++ code hase been compiled with VS13 and for python the anaconda distribution with Python 2.7 was used. Benchmarks were run on an idle computer with as little background processes running as possible. The sources of the benchmarks can be found on Github.

6.1.1 ModernGL

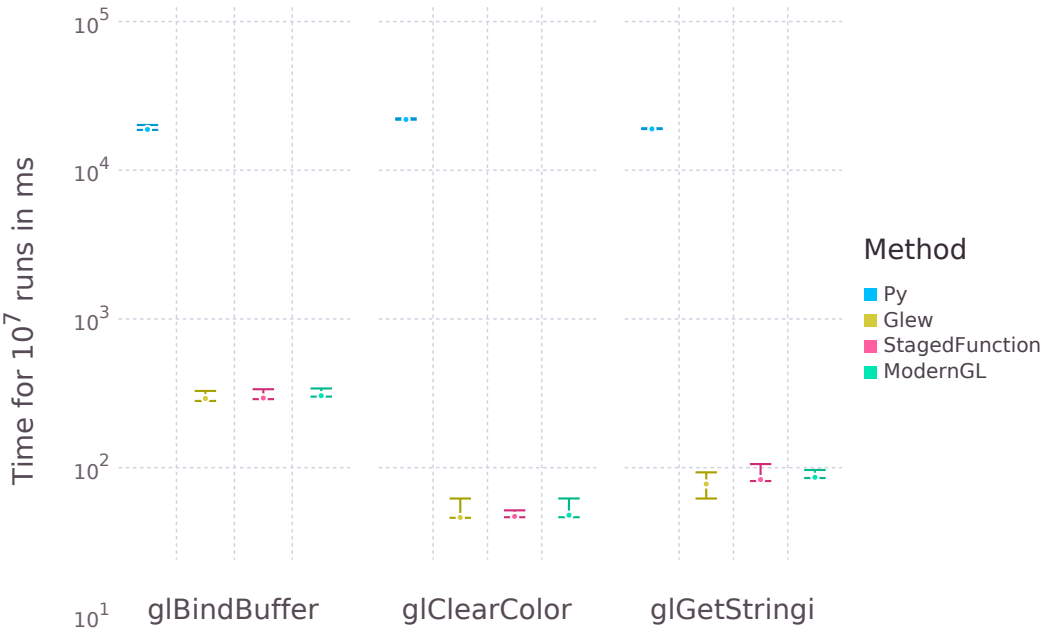

Figure 10: *Different performance of OpenGL wrappers. The time for* 10^7 *calls was measured 100 times for each function.*

Function	Python	Staged Function	ModernGL
glBindBuffer	64.43	1.00	1.04
glClearColor	474.72	1.02	1.04
glStringi	244.44	1.07	1.1

Table 3: *Performance relative to C++ with Glew (slowdown, bigger is worse)*

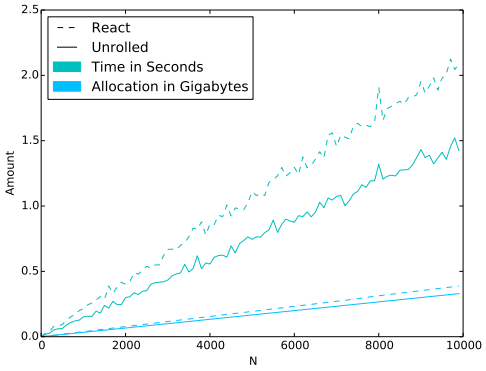

Figure 11: *Complicated graph, simple calculation*

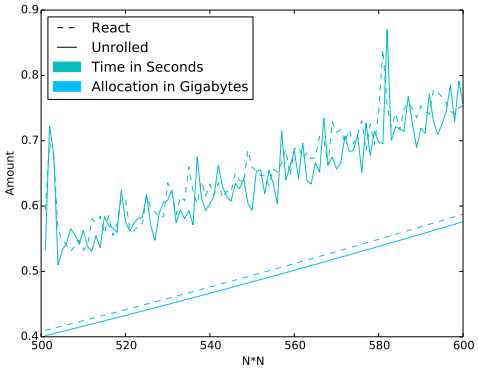

Figure 12: *High memory, simple event graph*

In this chapter, ModernGL, GLEW and PyOpenGL will get benchmarked. The procedure was, to call an OpenGL function 10^7 times in a tight loop. Execution time of the loop gets measured. The results are plotted in figure 10. ModernGL seems to do pretty well compared to C++ and python does very badly, with being up to 470 times slower in the case of glClearColor. Julia in contrast offers nearly the same speed as calling OpenGL functions from C++ as can be seen in the table 6. As all the OpenGL wrappers are pretty mature by now and bind to the same C library (the video driver), this should mainly be a C function call benchmark. Python performs badly here, but it must be noted that there are a lot of different Python distributions and some promise to have better C interoperability. As this benchmarks goal is to show that Julia's ccall interface is comparable to a C function call from inside C++, the python options have not been researched that thoroughly. From this benchmark can be concluded, that Julia offers a solid basis for an OpenGL wrapper library.

6.1.2 Reactive

It is relatively hard to benchmark the used event system in real world scenarious as it is hard to find a baseline. One would have to rewrite Romeo with another Event system. Using other visulization libraries as a baseline is also difficult, as it is hard to isolate the performance of the event system. This is why we will compare an event graph from Reactive with its unrolled version. For the unrolled version the functions from the callback graph have been executed in the same order as the event graph would have without introducing any event system related overhead. This way we can measure the overhead introduced by the event system. Two code samples have been benchmarked, one simulating the operation needed for the camera and the other simulates animating a large array. The first has low memory usage with a more complex event graph. The

second has a straight forward event graph, but it must pass on a large array and needs to execute the callbacks on the array.

As can be seen in figure 15, small operations with a complex event graph has some noticable overhead. Reactive is in this scenario about 1.45 times slower than the optimal version. This does not come as a surprise as sorting and managing the graph structure obviously adds some overhead.

The second scenario looks much better for Reactive. The performance difference is neglectable, making Reactive a good choice for creating signals with high memory throughput.

6.1.3 3D Rendering Benchmark

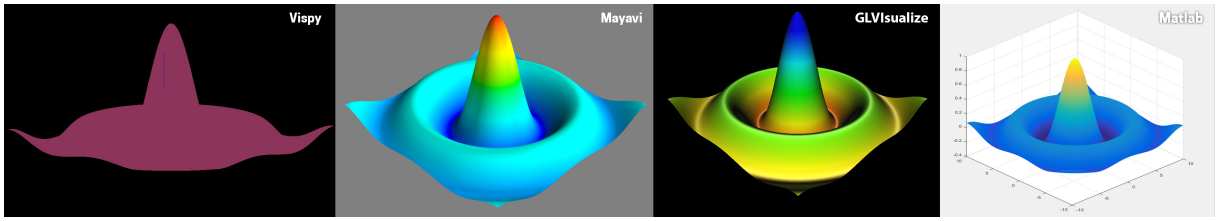

Figure 13: *Different visualizations of the same surface*

The biggest problem with benchmarking the 3D rendering speed is, that there is no library which will allow to exactly reproduce similar conditions and measures. Additionally, without extensive knowledge of the library, it is difficult to foresee what gets benchmarked. As an example of why it is difficult to measure the frame rate we can look at Vispy. When you enable to measure the frame rate, it will show very low frame rates, as it only creates a new frame on demand. On the other side Romeo has a fixed render loop, which renders as much frames as possible, leading to totally different amount of rendered frames per second. This is why it was decided, to use the threshold at which a similar 3D scene is still conceived as enjoyable and interactive. Usually the minimal amount of frames per second for perceiving movements as smooth is around 25. So the benchmark was executed in the way, that the number regulating the complexity of the 3D scene was increased until one could not move the camera without stutters. The recorded threshold is than the result of the Benchmark.

First benchmark is an animated and still 3D surface plot. The libraries offering this functionality where Vispy, Mayavi and Matlab.

Vispy had some issues, as the camera was never really smooth for the surface example. Also the normals were missing and there was no option to colorize the surface depending on the height. It was decided to use the threshold of going from a little stutter to unpleas-

Library	Still	Animated
Vispy	300	80
Mayavi	800	150
Matlab	800	450
Romeo	900	600
Speed up Vispy	9x	56x
Speed up Mayavi	1.26x	16x
Speed up Matlab	1.26x	1.7x

Table 4: *3D surface created from a NxN matrix.*

ant stutters, making vispy not completely fail this benchmark. For Vispy it was found out, that the normals where calculated on the CPU resulting in a major slow down[7]. The same can be expected for Mayavi, but Mayavi seems to be faster at calculating the normals. There is not much information available on how Matlab renders their visualization, as it is closed source.

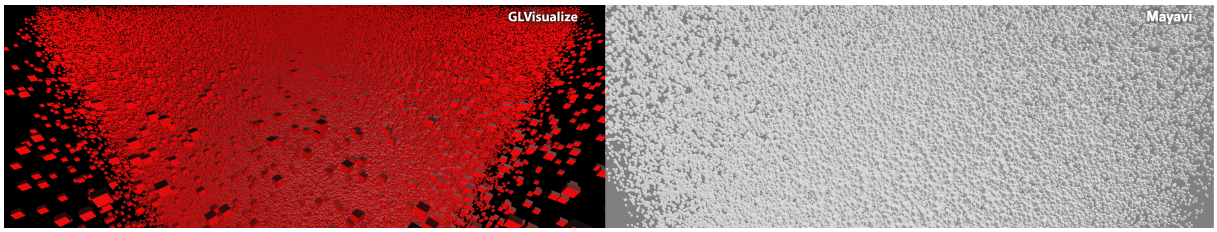

Figure 14: *Rendered particles*

The next benchmark is only between Romeo and Mayavi, as the other libraries did not offer a comparable solution. Matlab does not allow to use cubes as particle primitives and Vispy only had an example, where you needed to write your own shader, which can not be seen as a serious option. This is a benchmark for easy to use and high level plotting libraries. It is always possible to write an optimal version yourself in some framework, but what really interesting is, how well you can solve a problem with the tools the library has readily available.

Library	Still	Animated
Mayavi	90000	2500
Romeo	1000000	40000
Speed up	11x	16x

Table 5: *Maximum number of particles that could be displayed without stutter.*

Romeo is an order of magnitude faster in this specific benchmark. This is most likely due to the fact that Romeo uses OpenGL's native instance rendering.

6.1.4 IJulia

It was not possible to compare IJulia directly with Romeo, as the feature set for plotting is too different.

But there are certain factors, which indicate, that is hard to reach optimal performance with IJulia. First of all, IJulia uses ZMQ to bridge the web interface with the Julia kernel. ZMQ is a messaging system using different sockets for communication like inproc, IPC, TCP, TIPC and multicas. While it is very fast at it's task of sending messages, it can not compete with the native performance of staying inside one language. This is not very important as long as there does not have to be much communication between Julia and the IPython kernel. This changes drastically for animations, where big memory chunks have to be streamed to the rendering engine of the browser. It can be expected, that this will always be a weakness of IJulia. On the other hand, GPU accelerated rendering in a web browser is also limited. It relies on WebGL, which offers only a subset of the OpenGL's functionality. So while the execution speed of OpenGL can be expected to be similar, there are a lot of new techniques missing, which can speed up rendering.

To investigate this another benchmark has been created. It is between Romeo and Compse3D, which was the only library found to be able to display 3D models created with Julia directly from the IJulia notebook. This benchmark is not entirely fair, as Compose3D is just a very rough prototype so far. But there seems to be no other library with which you can easily create and display interactive 3D graphics in the IJulia or IPython notebook. This benchmark creates a sierpinsky gasket and Compose3D displays it in the IJulia notebook while Romeo displays it natively in a window.

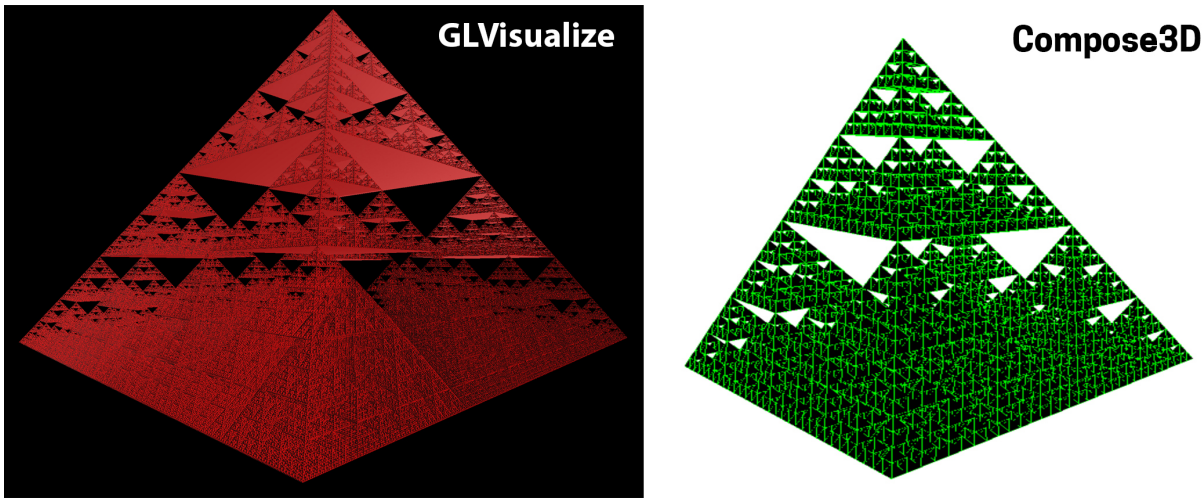

Figure 15: *Sierpinsky pyramid in 3D*

Again, Romeo is an order of magnitude faster. This can change in the future when Compose3D matures. But one needs to notice, that Romeo utilizes OpenGL's instancing

Library	Still
Compose3D	15625
Romeo	1953125
Speed up	125x

Table 6: *Maximum number of pyramids that could be displayed without stutter.*

to gain this speed. Native instancing is not yet available in WebGL, which means that this optimization will not be available for the IPython notebook in the near future.

6.2 Extensibility Analysis

The modular design of Romeo has proven to be effective and the goal of re-usability has already proven itself. Most of the created modules are used independently by different people. GLVisualize is used by myself for two packages, namely GLPlot, a scientific plotting package for Julia and for a prototype of a file explorer. It got forked by several users to create their own dynamic visualization packages. The same applies for ModernGL and GLAbstraction. Most other used packages are at least used by one other project. This indicates, that the abstraction and modularity is well designed, so that all the modules can function on their own.

The only exception is GLWindow, which has been used just indirectly through the other packages. This can mean three things. First, it is badly abstracted and does not cleanly wrap one use case. Secondly, it can be that the use case is not entirely clear to other people, which would not be a big surprise considering the minimal amount of documentation for GLWindow. And finally, considering the small group of people developing graphics for Julia, it could be that they simply do not need the lower level functionality of GLWindow and instead rely on the other written packages that use GLWindow.

Modularity guarantees a broad user and developer base, which in turn results in rich functionality and stability. From further analyzing the Github repository written for this thesis, one can find out that there is a general lack of documentation. This hinders people from contributing and using the packages.

The implementation in just one language has been achieved by choice. There are only a few exceptions, like the kernel code for OpenGL shaders, which currently can not be written in Julia. Julia programmers that use Romeo can extend Romeo with Julia and immediately see their results without complicated compilations. This together with the speed is one of the main achievements compared to other libraries offering similar functionality, like IJulia, Mayavi and Matlab. To further proof this point I will analyze the mentioned software in more detail. The language usage statistics and necessary tools needed in

order to extend the software will be the main focus of the analysis. One needs to note, that the usage statistic of languages is just a weak indicator for the extendability of a software. Using different languages for one project can make sense if the project has different domains where domain specific languages give an advantage. As this means one needs to know all used languages, this still introduces complexity, but it is at least justified complexity. This chapter will only discuss the complexity introduced by languages, which are only needed for compatibility with other libraries or because the main language is too slow. This is something, which should ideally be avoided.

6.2.1 IJulia

IJUlia is written in Julia and relies on ZMQ(C++) and IPython. IPython uses multiple JavaScript rendering back ends like Three.js and D3. [more to come]

Software	**languages used**
IPython	Python 78.5% JavaScript:15.1% HTML 5.0% Other 1.4%
Three.js:	JavaScript 62.4% HTML 26.4% Python 6.9% C++ 1.9% C 1.3% GLSL 0.6%
D3:	JavaScript 95.6% CSS 4.3%

Table 7: *]*
Technologies used in IJulia. Statistics taken from Github

6.2.2 Mayavi and VTK

Table 8 in the appendix shows an extensive summary of the used languages in the Paraview repository. It amounts to a total of 3.642.105 lines of code written in 29 languages. [more to come]

6.2.3 Matlab

Matlab is closed source, which makes the core of Matlab impossible to extend by the user. This is why Matlab relays on a plug-in architecture, which enables developers to write closed or open source plug-ins for Matlab. [Analysis of the plug-in Architecture] [more to come]

6.3 Usability Analysis

Doing a broad user survey or similar methods was out of scope for this thesis. As a result from this, the usability study has to be done analytically. There are different aspects which can be analyzed. For example, how many function names need to be remembered, how easy they are memorized, if they expose the wanted functionality and how difficult it is to look up unknown functionality. For this thesis, I will analyze the two main packages,

namely GLVisualizes and GLAbstraction. The named aspects will be analyzed and in addition feedback from Github will be used.

GLVisulize has a very simple API, as it offers only five functions: visualize, visualizedefaults, edit and renderloop. There are also the functions bounce and loop, which offers a simplification for creating periodic signals. These functions might get moved into Reactive, though. So for GLVisualize, only very few function names have to be remembered. The question is, does this simple interface still allow people to create the visualization they want. At closer inspection one can see, that visualize is overloaded 67 times, with each of these methods having a set of keyword arguments which enables further customization. These can introduce drastic changes. The particle visualization for example can take any mesh as a primitive. This enables a customization, which was not possible in such an easy way in the other examined packages. Also, most of the functions take either a data type, or a signal of that data type. This makes it very intuitive to animate your data. In contrast, in order to setup the animation for the other packages, it took quite some time to find out how to update values of an existing visualization. This is acceptable, as it might take quite some time to find out that Romeo uses signals and how to work with signals. But when this is found out, Romeo functions in the same, consistent way. Signals add a fourth dimension to any parameter or data you would like to visualize, making the usage principle consistent across the different visualizations.

For the other packages though, one needs to find out the names of the data for every visualization type in order to access and update them. Some attributes can not be animated, making the API even less consistent.

So for Romeo one can achieve anything by bringing the data into the right format. Problems arise, if this can not be done easily or the format is not intuitive for the programmer. To be fair, you will have this problem with every kind of visualization API. The difference is in the end, how easy it is to do the data transformations. Lets examine an example, where GLVisualize often will not allow to directly call visualize on the data. There is only a method for visualizing a Mesh, but not for a vertex list plus a face list. If you work with mesh data, you will often handle the face and vertex list isolated. So an API that offers a function like `visualize_mesh(x::VertexList, y::FaceList)` will be more straightforward for a programmer. Especially, as this is the standard way of displaying a mesh in most scientific plotting packages. This has become one of the most occurring questions on github, even though that there are usage examples for displaying a mesh.

But this functionality can be added in a simple way for GLVisualize by defining `visualize_mesh(face vertexlist) = visualize(Mesh(facelist, vertexlist))`. As GLVisulize should stay

as close to the principle of only having one function name, this should be moved to other interface only packages, though. [more to come]

7 Conclusion

7.1 Future Work

8 References

[1] Open Benchmarking. Llvm clang 3.6 compiler tests. Accessed: 04/15/2015 : http://www.webcitation.org/6XoVc1S0y.

[2] J. Bezanson, S. Karpinski, V. B. Shah, and A. Edelman. Julia: A Fast Dynamic Language for Technical Computing. *ArXiv e-prints*, September 2012.

[3] Evan Czaplicki. Elm. Accessed: 05/22/2015 : http://www.webcitation.org/6YiVPsavP.

[4] Al Danial. Cloc. Accessed: 04/01/2015 : http://www.webcitation.org/6XT73jFkv.

[5] Enthought. Mayavi. Accessed: 05/16/2015 : http://www.webcitation.org/6YZNh5oEC.

[6] Enthought. Vispy. Accessed: 05/16/2015 : http://www.webcitation.org/6YZPFZODy.

[7] Github. surface slow and buggy 892. Accessed: 05/13/2015 : http://www.webcitation.org/6YVE1zeNb.

[8] GLFW. Glfw. Accessed: 05/22/2015 : http://www.webcitation.org/6YiVA6PY7.

[9] Sebastion Good. Little performance explorations. Accessed: 04/14/2015 : http://www.webcitation.org/6Xmtrox4u.

[10] Shashi Gowda. Reactive. Accessed: 05/22/2015 : http://www.webcitation.org/6YiVI2SUy.

[11] Chris Green. Improved alpha-tested magnification for vector textures and special effects. In *ACM SIGGRAPH 2007 CoWurses*, SIGGRAPH '07, pages 9–18, New York, NY, USA, 2007. ACM.

[12] IPython. Ijulia notebook. Accessed: 03/23/2015, Archived by WebCite®: http://www.webcitation.org/6XFFIHQee.

[13] Viral Shah Alan Edelman Jeff Bezanson, Stefan Karpinski. Why we created julia. Accessed: 04/01/2015 : http://www.webcitation.org/6XT6wqYAf.

[14] Kitware. Vtk gallery. Accessed: 04/15/2015 : http://www.webcitation.org/6XodgQgdm.

[15] Julia Lang. Calling c and fortran code. Accessed: 04/02/2015 : http://www.webcitation.org/6XUueZYLw.

[16] Julia Language. benchmark times. Accessed: 04/13/2015 : http://www.webcitation.org/6XlX8Wwft.

[17] Michael Larabel. Compiler intel broadwell linux tests. Accessed: 04/15/2015 : http://www.webcitation.org/6XoVX2L1r.

[18] Michael Larabel. Intel broadwell: Gcc 4.9 vs. llvm clang 3.5 compiler benchmarks. Accessed: 04/15/2015 : http://www.webcitation.org/6XoW1HeDq.

[19] Michael Larabel. Microsoft announces an llvm-based compiler for .net. Accessed: 04/15/2015 : http://www.webcitation.org/6YOdC964v.

[20] Eivind Lyngsnes Liland. Path rasterizer for openvg, 2007.

[21] MathWorks. Matlab pricing. Accessed: 03/22/2015, Archived by WebCite®: http://www.webcitation.org/6XEFJ0PBE.

[22] Microsoft. Wglgetprodaddress documentation. Accessed: 02/27/2015, Archived by WebCite®: http://www.webcitation.org/6WemKehYL.

[23] Tanmay Mohapatra. reduce gc load in readdlm. Accessed: 04/15/2015 : http://www.webcitation.org/6YOc9Qv1T.

[24] Amuthan Arunkumar Ramabathiran. Finite element programming in julia. Accessed: 04/14/2015 : http://www.webcitation.org/6XmvHthh5.

[25] Tracy Wadleigh. Iso surfaces. Accessed: 05/24/2015 : http://www.webcitation.org/6Yltap0ze.

[26] OpenGL Wiki. Rendering pipeline overview. Accessed: 04/10/2015 : http://www.webcitation.org/6Xgd8fedi.

[27] Wikipedia. Ipyton. Accessed: 03/23/2015, Archived by WebCite®: http://www.webcitation.org/6XFEB9BB3.

[28] Wired. The one last thread holding apple and google together. Accessed: 04/15/2015 : http://www.webcitation.org/6YOdawS5T.

Appendix

A IJulia

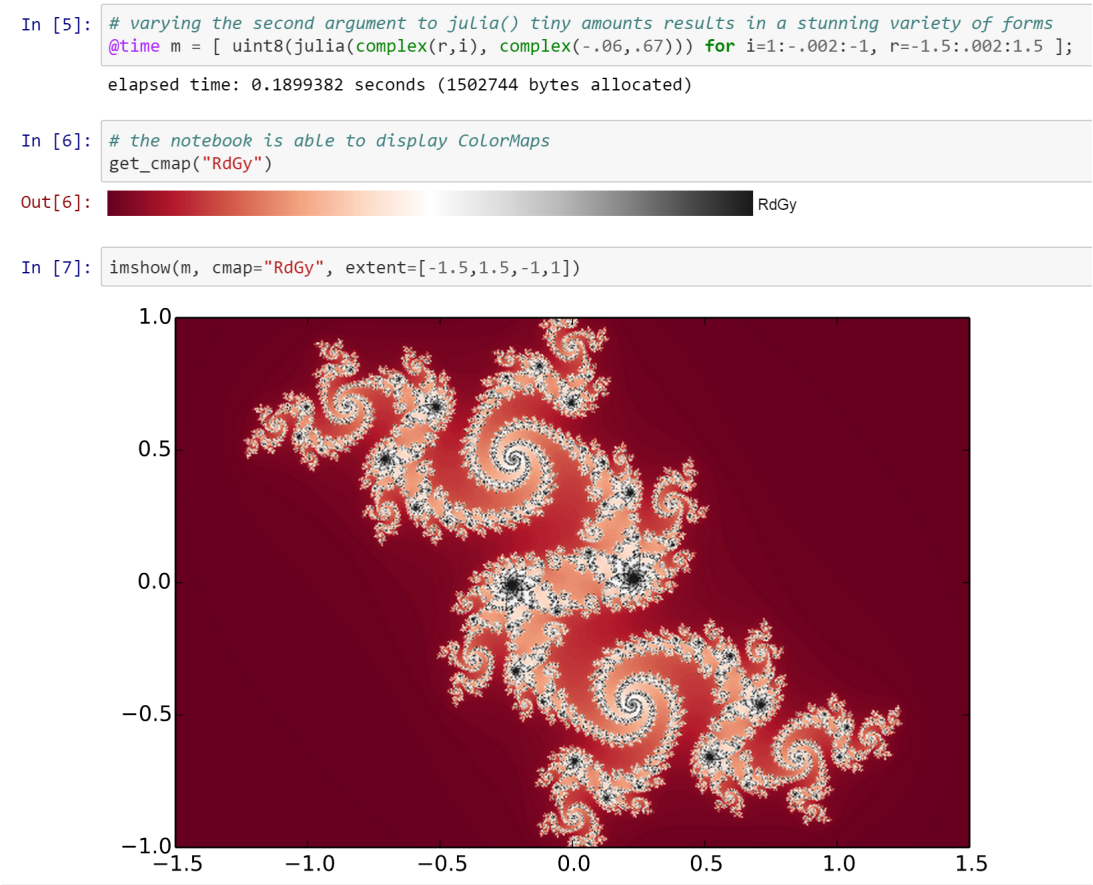

Figure 16: *Example of an IJulia Notebooks. Screenshot taken from [12]*

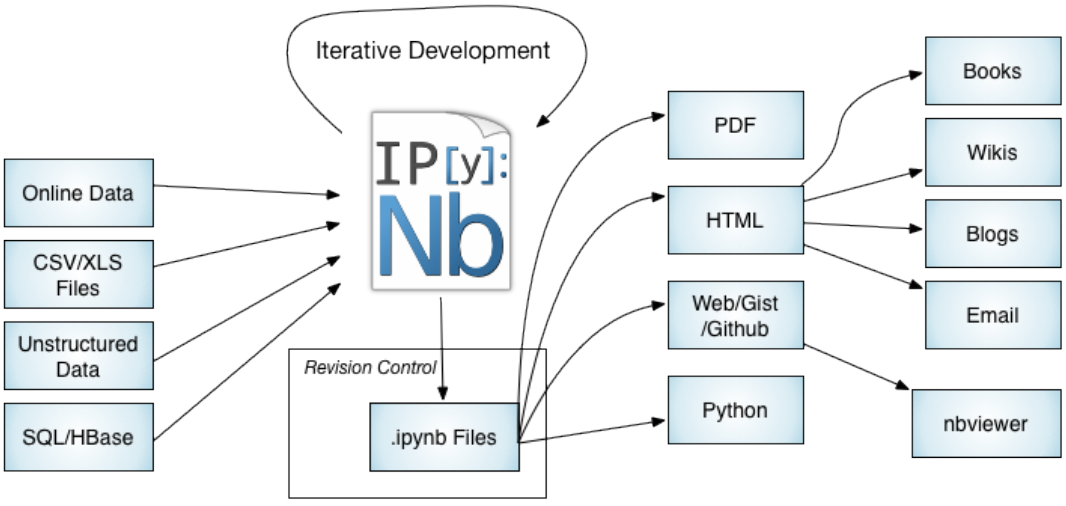

Figure 17: *Workflow of IPython Notebooks. Graphic from Wikipedia [27]*

B Language Statistics

All language statistics have been made with cloc [4] the current master of the github repositories.

Table 8: *Paraview, language statistic*

Language	**files**	**blank**	**comments**	**code**
C++	2037	70003	86594	391121
C/C++ Header	1937	48345	93434	141581
C	273	35843	17101	135937
XML	275	1930	3521	59030
Fortran 77	67	28	18039	39116
Python	209	5883	8719	21935
CMake	443	3705	6185	20025
Javascript	20	1285	1847	7982
CSS	23	750	251	4827
HTML	26	240	1692	2328
JSON	13	2	0	2162
yacc	1	207	138	881
Bourne Again Shell	19	186	347	799
make	8	248	90	734
Bourne Shell	18	158	116	708
XSLT	3	46	17	388
CUDA	1	58	184	318
Pascal	2	69	102	228
SUM:	5375	168986	238377	830100

Table 10: *VTK, language statistic*

Language	files	blank	comment	code
C++	3845	203851	179827	1278279
C	1103	130996	289623	707122
C/C++ Header	3489	103162	246368	382728
Python	1681	88983	121122	258787
Tcl/Tk	573	11052	7830	48213
CMake	739	4715	7424	35956
Javascript	47	6941	6747	33098
CSS	33	1476	323	18100
XML	10	17	36	8337
Objective C++	20	1210	1372	5601
m4	3	660	83	4922
yacc	3	726	570	4852
HTML	25	553	531	4313
Java	50	912	1192	4239
Cython	20	848	1625	3484
Perl	11	939	950	3119
JSON	3	5	0	2658
Windows Resource File	21	333	380	1835
lex	3	215	162	1510
DTD	3	435	477	1335
Assembly	13	202	0	936
Bourne Again Shell	16	191	333	866
CUDA	6	113	77	740
Bourne Shell	15	64	122	380
make	5	54	187	170
IDL	1	0	0	150
Windows Module Definition	3	3	0	142
JavaServer Faces	3	26	0	88
Objective C	2	13	18	17
SUM:	11749	558698	867379	2812005

C Romeo's GUI

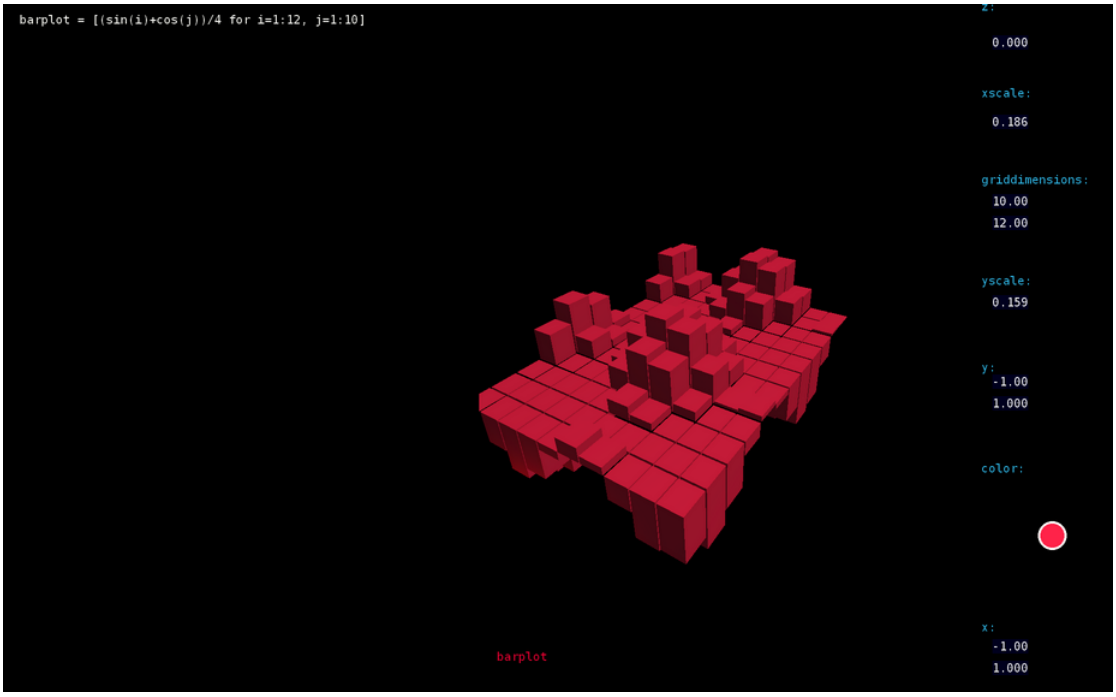

Figure 18: *Screenshot of the prototype. Left: evaluated script, middle: visualization of the variable barplot, right: GUI for editing the parameters of the visualization*

D Benchmark

implementations	Language	Speed in Seconds
JFinEALE	Julia	9.6
Comsol 4.4 with PARDISO	Java	16
Comsol 4.4 with MUMPS	Java	22
Comsol 4.4 with SPOOLES	Java	37
FinEALE	Matlab	810

Table 12: *FE Implementation comparison*

Benchmark	LLVM35	LLVM36	Difference
Rodinia	265.34	289.43	0.92
Rodinia	118.52	118.42	1.00
FFTW	6034.26	5961.52	0.99
FFTW	5988.02	5969.68	1.00
FFTW	4373.76	4349.26	0.99
FFTW	4405.26	4376.60	0.99
Timed HMMer Search	11.43	12.20	0.94
Timed MAFFT Alignment	4.69	4.69	1.00
SciMark	2008.04	1939.20	0.97
SciMark	556.37	537.07	0.97
SciMark	355.18	362.56	1.02
SciMark	2790.01	2452.42	0.88
SciMark	4843.14	4834.33	1.00
SciMark	1495.53	1509.64	1.01
John The Ripper	936.00	984.00	1.05
John The Ripper	5219000.00	5204000.00	1.00
John The Ripper	14767.00	14779.00	1.00
Himeno Benchmark	1572.74	1574.91	1.00
Timed Apache Compilation	23.56	25.34	0.93
Timed ImageMagick Compilation	19.13	20.67	0.93
C-Ray	12.13	12.73	0.95
Smallpt	148.00	145.00	1.02
Stockfish	3775.00	3812.00	0.99
Bullet Physics Engine	3.43	3.40	1.01
Bullet Physics Engine	5.85	5.95	0.98
Bullet Physics Engine	6.38	6.51	0.98
Bullet Physics Engine	5.99	5.68	1.05
Bullet Physics Engine	3.77	3.84	0.98
Bullet Physics Engine	1.25	1.23	1.02
Bullet Physics Engine	1.47	1.46	1.01
FLAC Audio Encoding	7.17	7.28	0.98
LAME MP3 Encoding	15.99	15.43	1.04
Hierarchical INTegration	240423016.30	264346632.10	1.10
Apache Benchmark	17643.10	19412.76	1.10

Table 14: *LLVM 3.5 compared to LLVM 3.6 in the Phoronix benchmark test suite[1]*

Benchmark	GCC492	LLVM35	Difference
Timed MAFFT Alignment	11.39	12.88	0.88
Timed MrBayes Analysis	25.44	26.28	0.97
SciMark	1179.35	1497.26	1.27
SciMark	564.44	603.62	1.07
SciMark	263.97	279.26	1.06
SciMark	1957.09	2070.94	1.06
SciMark	2046.08	2953.00	1.44
SciMark	1065.17	1579.54	1.48
John The Ripper	2382.00	926.00	0.39
John The Ripper	4296333.00	4925333.00	1.15
Himeno Benchmark	1618.11	1459.12	0.90
ebizzy	18192.00	17879.00	0.98
Timed Apache Compilation	55.65	38.61	1.44
Timed PHP Compilation	60.94	44.14	1.38
C-Ray	48.40	73.93	0.65
Smallpt	64.00	148.00	0.43
Stockfish	3933.00	4108.00	0.96
Bullet Physics Engine	5.75	6.08	0.95
Bullet Physics Engine	6.65	7.45	0.89
Bullet Physics Engine	5.76	6.59	0.87
Bullet Physics Engine	3.79	3.89	0.97
Bullet Physics Engine	1.23	1.31	0.94
Bullet Physics Engine	1.46	1.57	0.93
FLAC Audio Encoding	6.87	9.12	0.75
LAME MP3 Encoding	12.82	12.88	1.00
Hierarchical INTegration	206580965.69	237608504.18	1.15
Apache Benchmark	15429.86	15499.88	1.00
FFTW	6283.32	5443.12	0.87
FFTW	6053.00	5447.48	0.90
FFTW	4504.06	4260.74	0.95
FFTW	4091.20	4070.66	0.99

Table 16: *gcc 4.9.2 compared to LLVM 3.5 in the Phoronix benchmark test suite[17]*

Official Statement

I hereby guarantee, that I wrote this thesis and didn't use any other sources and utilities than mentioned.

Date: .....

(Signature)